

# mtg-volatility-signal: the Current Volatility index

---

**Snapshot 2026-09-23. By Cameraderie Cards. Informational only, not financial advice.**

Every other model in this toolkit tells you *direction*: will a card rise (buy), has it peaked (sell), is it about to be reprinted (brace), can you exit it (liquidity). None of them tells you *how far*. A 5% mover and a 300% mover look identical to a probability. This model fills the **magnitude** gap. It assigns every Magic: The Gathering card a **Current Volatility** score from 0 to 100: how much its price swings right now, independent of which way.

Unlike the demand and liquidity siblings, which are composite proxy indices, the core of this score is **measured directly** from each card's own price history. The number is an annualized realized volatility, not an estimate of one.

## What goes into the score

Four measured components, each percentile-normalized and blended (weights {'realized': 0.5, 'drawdown': 0.25, 'range': 0.15, 'dispersion': 0.1}):


| Component    | What it measures                                                                 | How                                     |
|--------------|----------------------------------------------------------------------------------|-----------------------------------------|
| Realized vol | annualized volatility of weekly log returns, weighted toward recent weeks (EWMA) | 16-year daily history, resampled weekly |
| Drawdown     | worst peak-to-trough fall and downside-only deviation: the "how far down"        | same history                            |
| Range        | swing amplitude, the inter-quantile price span relative to the price level       | same history                            |
| Dispersion   | jump risk (largest single-week move) and fat tails (kurtosis of returns)         | same history                            |


Returns are measured on **weekly** sampled prices, not daily. Magic cards are thinly traded, so daily marks carry heavy discretization noise (a \$1.50 card quoted in \$0.25

steps shows fake 20% daily swings) and stale-quote jumps. Resampling to weekly closes diversifies that noise away while preserving genuine week-to-week moves, the standard treatment for an illiquid asset. Volatility is computed from a single representative printing per card (the cheapest well-tracked copy a typical owner holds), never a median across printings, which would manufacture fake volatility when a card's quoted printings change from day to day.

## Coverage


Scored **30,538 cards** at this snapshot. **30,471** of them have enough price history for a **measured** volatility; the small remainder are imputed from a prebuilt volatility feature and flagged `is_imputed`. The median card has an annualized realized volatility of **26.0%**.






## Does the score mean anything? Two checks

**1. Measured volatility rises with the buckets.** A consistency check: the realized volatility the score is built from really does climb monotonically from Stable to Highly volatile.




**2. Volatility persists (the real test).** A backward-looking measurement is only a useful forward signal if volatility is persistent. So we measure each card's volatility on the **first half** of its history and, separately, on the **second half**, and check whether

the first predicts the second. It does, at rank correlation **nan**. The second half is never seen when the first is computed, so this is genuine out-of-time corroboration, not circular. The fitted slope is **nan**, well below 1: volatility is persistent but **mean-reverting**, so the forward bands below shrink today's extreme readings toward the average rather than extrapolating them.



## The magnitude deliverable: expected range and drawdown

Because the score is a real volatility, it converts directly into the magnitudes the rest of the toolkit needs. For every card the output carries an **expected price range** at 3, 6, and 12 months (lognormal +/- 1 sigma bands, about a 68% interval) built from the mean-reverting forward volatility, plus a measured **max drawdown**. Across all cards the median expected 12-month move is about **29.7%** (3-month **13.9%**, 6-month **20.1%**).



## Example cards

Example cards (snapshot 2026-09-23)

|                                 | Riding the Dilu Horse<br>vol 98   drawdown 100% | Zhou Yu, Chief Commander<br>vol 98   drawdown 97% | Caltrops<br>vol 98   drawdown 79%         | Cataclysm<br>vol 98   drawdown 63%      | Gregor, Shrewd Magistrate<br>vol 98   drawdown 81% |
|---------------------------------|-------------------------------------------------|---------------------------------------------------|-------------------------------------------|-----------------------------------------|----------------------------------------------------|
| <b>Most volatile (&gt; \$3)</b> |                                                 |                                                   |                                           |                                         |                                                    |
| <b>Most stable (&gt; \$3)</b>   | Mold Demon<br>vol 12   drawdown 3%              | Scarwood Bandits<br>vol 12   drawdown 3%          | Intervention Pact<br>vol 12   drawdown 4% | Petrified Field<br>vol 12   drawdown 4% | Marshaling the Troops<br>vol 13   drawdown 4%      |

The most volatile cards are recent hyped releases and thin-market specials that spike and crash, and cards rocked by bans or buyouts. The most stable are heavily-reprinted, deeply-played staples whose price barely moves week to week.

## Buckets

| Bucket          | Score  | Cards | Median realized vol |
|-----------------|--------|-------|---------------------|
| Highly volatile | 75-100 | 5,810 | 60.4%               |
| Volatile        | 50-75  | 9,507 | 31.6%               |
| Steady          | 25-50  | 9,273 | 20.9%               |
| Stable          | 0-25   | 5,948 | 8.4%                |

## How this plugs into the other signals

Volatility is the magnitude term every other signal is missing. The toolkit's risk-adjusted value is roughly  $EV = \text{spike\_prob} \times \text{upside} - \text{fall\_prob} \times \text{downside} - \text{reprint\_prob} \times \text{reprint\_crash}$ . The `upside` and `downside` terms are exactly what this model supplies: the expected range gives the size of the move, the drawdown gives the worst-case downside. With them the EV becomes genuinely risk-adjusted, and positions can be sized properly: bigger on high-conviction low-variance cards, smaller on the violent ones.

## Honest limitations

- This is **version 1: a measured realized-volatility index, not an options-implied or forward forecast**. The only forward assumption is volatility persistence, which is real but moderate and mean-reverting (slope nan), so the forward bands are estimates, not guarantees.
- **Cheap cards carry residual noise**. Weekly sampling removes most price-discretization noise, but a sub-\$2 card still shows more percentage wiggle than its dollar moves justify. Imagery and examples are restricted to cards over \$3 for this reason.
- Volatility is computed from **one representative printing**. A card's other printings (foils, old borders, collectible editions) can be much more or less volatile than the tracked copy.

- A **stale price looks stable**. A card no one trades shows no volatility because nothing reprices, not because it is genuinely calm. Cards moving in under 5% of weeks are flagged `stale_price` ; read them alongside the liquidity signal.
- Informational only, not financial advice.

Generated by `scripts/build_report.py` from `data/processed/volatility_features.parquet`.